

CS6350.002 Final Project Report

Detection and Localization of Brain Tumor from MRI Scan Images. *

Mahsa Eslamialishah*, Pritom Das Radheshyam† and Novonil Das‡

The University of Texas at Dallas

*mxe190002@utdallas.edu, †pxr180033@utdallas.edu, ‡nxd180022@utdallas.edu

Abstract—In recent years there has been growing popularity of Deep Learning in almost every fields where decision making is involved - finance, health care, marketing, sales and what not. Deep Learning has shown promising results in the field of healthcare in many areas such as: Disease Diagnosis with Medical Imaging, Surgical Robots, Maximizing Hospital Efficiency. AI healthcare market is expected to reach \$45.2 billion USD by 2026 from the current valuation of \$4.9 billion USD. Deep learning has been proven to be superior in detecting diseases from X-rays, MRI scans and CT scans which could significantly improve the speed and accuracy of diagnosis.

In this CS6350.002 final project we have applied Deep Learning to detect brain tumours from MRI Scan images using Residual Network and Convolutional Neural Networks. This automatic detection of brain tumors can improve the speed and accuracy of detecting and localizing brain tumors based on MRI scans. This would drastically reduce the cost of cancer diagnosis and help in early detection of tumors without any human involvement and would essentially be a life saver. We have also compared the accuracy of results obtained by using two different models - ResNet50 and ResNet18 and used Transfer Learning to tune or freeze weights to evaluate what gives best result.

We have 3929 brain MRI scans which are either positive or negative cases of brain tumour. We have built models using ResNet50 and ResNet18 and evaluated their performance in detecting positive or negative cases of brain tumors.

Index Terms—Brain Tumor Detection, Deep Learning, Residual Network, Convolutional Neural Network

I. INTRODUCTION

II. RELATED WORK

A. Image Segmentation

Image Segmentation is a process of classifying every pixel in an image into a designated category. Image Segmentation allows us to understand and extract information from images at the pixel-level. Image Segmentation can be used for object recognition and localization which offers tremendous value in many applications such as medical imaging and self-driving cars etc. We could use image segmentation to train a neural network to produce pixel-wise mask of the image. Modern image segmentation techniques are based in deep learning approach which makes use of common architectures such as CNN, FCNs(Fully Convolution Networks) and Deep Encoders-Decoders.

B. Convolutional Neural Network

Convolutional Neural Network (CNN) is a type of deep neural network which has proven to perform well in computer vision tasks such as image classification, object detection, object localization and neural style transfer. CNN extracts features from images, which an algorithm like MLP or RNN can't achieve. The architecture of a convolutional neural network looks something like this:

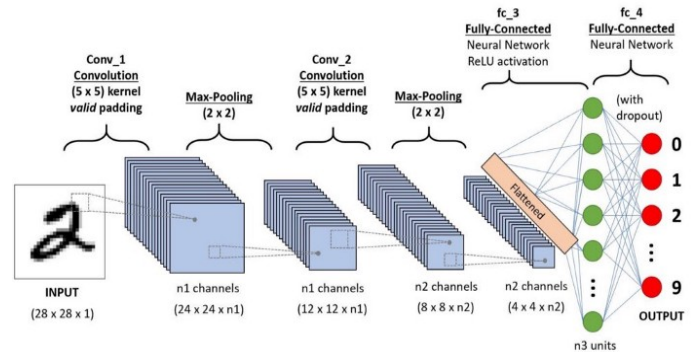


Fig. 1. Convolutional Neural Network

We should note a couple of things from this. The network starts off with 2 convolutional and max-pooling layers, followed with 2 fully connected layers which end with an output layer.

C. Residual Network

A residual neural network (ResNet) is a kind of artificial neural network which builds on constructs that are known as pyramidal cells found in the cerebral cortex. Residual neural networks do this by making use of skipping connections, or shortcuts to jump over some layers. Usually ResNet models are constructed with multiple layer skips that contain nonlinearities (ReLU) and can also have batch normalization implemented between them.

Although traditionally more layers mean a better network but because of the vanishing gradient problem, weights of the first layer won't be updated correctly through the back-propagation. Since back-propagation of the error gradient is done to earlier layers, repeated multiplication makes this




Fig. 2. Residual Network

gradient to become really small. Thus, problem can be taken care of by Res-Net solves this problem by using the identity matrix. When the back-propagation is done through identity function, the gradient will be multiplied only by 1 which helps in preserving the input and avoids any loss of information.

The network component consists of 3X3 filters, CNN down-sampling layers with stride 2, global average pooling layer and a 1000-way fully-connected layer with softmax in the end. ResNet helps in avoiding connection in which an original input is also added to the output of the convolution block which greatly helps to solve the problem of vanishing gradient by allowing an alternative path for the flow of gradient. Another big advantage is that it identifies functions that helps higher layer to perform as good as a lower layer. Traditionally each layer feeds into the next layer. But since this model's each layer feeds into the next layer and directly into the layers about some hops away.

D. Transfer Learning

A major assumption in many machine learning and data mining algorithms is that the training and future data must be in the same feature space and have the same distribution. However, in many real-world applications, this assumption may not hold. For example, we sometimes have a classification task in one domain of interest, but we only have sufficient training data in another domain of interest, where the latter data may be in a different feature space or follow a different data distribution. In such cases, knowledge transfer, if done successfully, would greatly improve the performance of learning by avoiding much expensive data-labeling efforts. Humans have an inherent ability to transfer knowledge across tasks. What we acquire as knowledge while learning about one task, we utilize in the same way to solve related tasks. Transfer Learning is a machine Learning technique in which a network that has been trained to perform specific task is being reused as a starting point of another similar task. Transfer Learning is widely used since starting from a pre-trained model can dramatically reduce the computational time required if training is performed from scratch [6]. The more related the tasks, the easier it is for us to transfer, or cross-utilize our knowledge. Human don't learn everything from

ground but we transfer and leverage our knowledge from the past! Conventional machine learning and deep learning algorithms, so far, have been traditionally designed to work in isolation. These algorithms are trained to solve specific tasks. Transfer learning is a kind of idea where we overcome the isolated learning paradigm and it utilizes knowledge acquired from a completed task. Traditional way of machine learning and deep learning approach thus far have been implemented to work in an isolated manner and solve only specific kind of tasks. Every time the models needed to be rebuilt from scratch once the feature-space distribution changes. Transfer learning is a novel approach which helps to overcome the issues of overcoming isolated learning paradigm and utilizing knowledge acquired for one task to solve related ones [1].

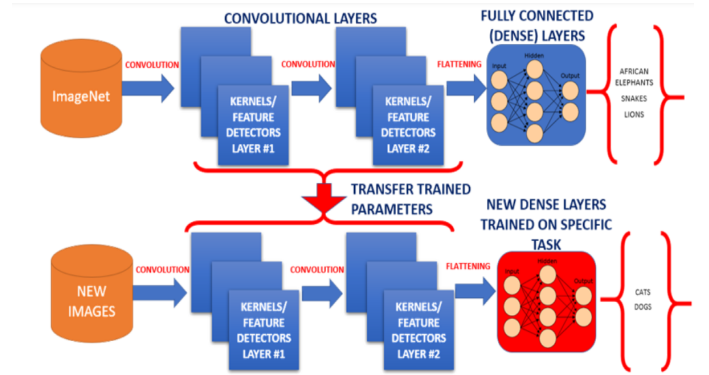


Fig. 3. Transfer Learning

One of the major use case of deep learning is the fact that most models solve very difficult problems and need a whole lot of data. Since it can features of it's own, thus it eliminates problems of labeling data that is usually done for supervised learning problems in machine learning.

III. DATA FORMATTING

We have worked on 3929 Brain MRI Scans. These are image files that contain the actual data of patients brain MRI scans and the information if they were analyzed by the doctor to have tumors or not. The data was converted into .png format and we ran the machine learning algorithm in all those images. We have built two different models - one using ResNet50 and the other using ResNet18, and also tuned certain hyper parameters to assess and compare the results of all the approaches. We used 85% of the data to train the model and the rest 15% of the data was used for testing.

IV. PROPOSED METHOD

We used Brain MRI Scans images as dataset to train our machine and test our machine learning model. We used multiple features to run our model and using Residual Networks we classified if the MRI scans would be classified to have tumors or not. We used 2 different models (ResNets) - ResNet50 and ResNet18 to classify the MRI scans to either contain tumors or not and tuned multiple hyper parameters in transfer learning - a. Freeze the trained CNN network weights from the first

layers, only train the newly added dense layers(with randomly initialized weights). b. Initialize the CNN network with the pre-trained weights and retrain the entire CNN network while setting the learning rate to be very small, this is critical to ensure that you do not aggressively change the trained weights.

A. Image Segmentation

The image is converted into a vector and possibly a classification head is added at the end. Classical convolutional neural networks are generally used when the entire image needs to be classified as a class label. Softmax function is applied to every pixel which makes the segmentation problem work as a classification problem where classification is performed on every pixel of the image.

B. Mask

The goal of the image segmentation is to understand the image at the pixel level. It associates each pixel with a certain class. The output produced by image segmentation model is called a "mask" of the image.

Masks can be represented by associating pixel values with their coordinates. For example if we have a black image of shape (2,2), this can be represented as $[[0,0],[0,0]]$. If the output mask is as follows $[[255,0],[0,255]]$. To represent this mask we have to first flatten the image into a 1-D array. This would result in something like $[255,0,0,255]$ for the mask. The, we can use the index to create the mask. Finally we would have something like $[1,0,0,1]$ as our mask.

C. Convolutional Neural Networks

Neural networks are a powerful technology for classification of visual inputs arising from documents. The first CNN layers are used to extract high level general features. The last couple of layers are used to perform classification (on a specific task). Local receptive fields scan the image first searching for simple shapes such as edges/lines. These edges are then picked up by the subsequent layer to form more complex features.

D. ResNet(Residual Neural Network)

Convolutional Neural Network (CNN) based models have achieved great success in Single Image Super-Resolution (SISR). Owing to the strength of deep networks, these CNN models learn an effective nonlinear mapping from the low-resolution input image to the high-resolution target image, at the cost of requiring enormous parameters. As convolutional neural networks grow deeper, vanishing gradient tend to occur which negatively impact network performance. Vanishing gradient descent problem occurs when the gradient is back-propagated to earlier layers which result in a very small gradient. Residual Neural Network includes "skip connection" feature which enables training of 152 layers without vanishing gradient issues. ResNet works by adding "identity mappings" on top of the CNN. ImageNet contains 11 million images and 11000 categories. ImageNet is used to train ResNet deep network.

E. ResNet50

ResNet-50 is a convolutional neural network that is 50 layers deep. You can load a pretrained version of the network trained on more than a million images from the ImageNet database [1]. The pretrained network can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. As a result, the network has learned rich feature representations for a wide range of images. The network has an image input size of 224-by-224. This can be used to classify new images using the ResNet-50 model.




Fig. 4. ResNet50

F. ResNet18

ResNet-18 is a convolutional neural network that is 18 layers deep. You can load a pretrained version of the network trained on more than a million images from the ImageNet database [1]. The pretrained network can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. As a result, the network has learned rich feature representations for a wide range of images. The network has an image input size of 224-by-224. This can be used to classify new images using the ResNet-50 model.

G. Transfer Learning

Transfer Learning is a machine learning technique in which a network that has been trained to perform a specific task is being reused or re-purposed as a starting point for another similar task. Transfer Learning is widely used since starting from a pre-trained model can dramatically reduce the computational time required if training is performed from scratch [6].

1) Transfer Learning Training Strategies:

- Strategy 1 Steps:
 - Freeze the trained CNN network weights from the first layers.
 - Only train the newly added dense layers(with randomly initialized weights).

[6]

- Strategy 2 Steps:

- Initialize the CNN network with the pre-trained weights.
- Retrain the entire CNN network while setting the learning rate to be very small, this is critical to ensure that you do not aggressively change the trained weights.

[6]

- Transfer Learning Advantages are:
 - Provides fast training progress, we don't need to start from the beginning using randomly initialized weights.
 - We can use small training data set to achieve incredible results.

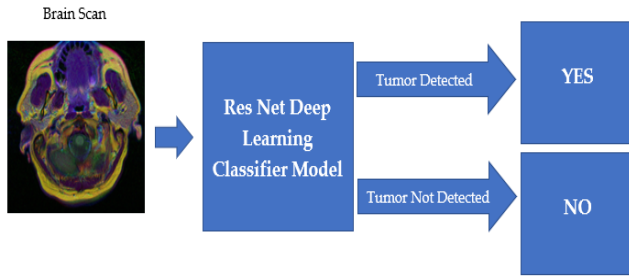


Fig. 5. Architecture

V. EXPERIMENTS

In this study, we explore different approaches for data augmentation and transfer learning. For data augmentation, we have leveraged random horizontal flipping on the training data. For transfer learning, we have investigated leveraging the pre-training models as feature extractors and an initial model. In the feature extractor scenario, the pre-trained layers are frozen during training the new task and only the added classification layers are trained. In the initialization scenario, the entire model, including the ResNet layers, is trained to learn the new task.

Since our training data is small compared to the number of trainable parameters, we hypothesized that a small pre-trained model would perform better for the initialization scenario. Therefore, we examine the ResNet50 and ResNet18 as our pre-trained models.

A. Experiment 1

The images files contain brain MRI scans and those are used as inputs to the model to train and test the model and in this experiment the images are resized.

B. Experiment 2

The images files contain brain MRI scans and those are used as inputs to the model to train and test the model and in this experiment the images are resized and normalized.

C. Experiment 3

The images files contain brain MRI scans and those are used as inputs to the model to train and test the model and in this experiment the images are resized and random horizontal flip.

D. Experiment 4

The images files contain brain MRI scans and those are used as inputs to the model to train and test the model and in this experiment the images are resized and random horizontal flip and normalized.

E. Experiment 5

VI. RESULT

We conducted a series of experiments and the results of those experiments have been attached here at the end. The results were then compared, and we found the following - By comparing experiment 1 and 3, where we changed the model from ResNet50 to ResNet18, we show that we achieved 1% improvement. this improvement is because ResNet18 has less number of parameters and since our data is small, ResNet50 gets overfitted to the data. By tuning the Neural Network hyperparameters we achieved 96% accuracy for ResNet18 and 94% accuracy for ResNet50. In the next experiments we have fixed these hyperparameters and normalized our data with mean = [0.485,0.456,0.406] and std= [0.229,0.224,0.225]. we show that our accuracy dropped from 96% to 92% for ResNet18 and improved by 1% for ResNet50. comparing experiment 13 and 15 shows that by adding random horizontal flipping and removing the normalization, the ResNet18 achieves a better performance. For the last experiment we have added random horizontal flipping and normalization. As the result of experiment 19 and 20 shows, we have achieved our best performance for both ResNet50 and ResNet18. We achieved the best results for a train batch size of 8 and test batch size of 8 with 30 epochs and Learning rate of 0.0001 and momentum of 0.5, with added layers of Linear RELU Dropout Linear Dropout Linear, and a transform on the dataset with a resize factor of 224 and Random Horizontal Flip we achieved an accuracy of 97% using a ResNet18 model and an accuracy of 96% using ResNet50 model.

VII. RUNNING PYSPARK JOBS IN AWS EMR CLUSTER VS AWS EMR NOTEBOOK

A. AWS EMR Notebook

One can use Amazon EMR Notebooks along with Amazon EMR clusters. It runs behind Apache Spark to create and open Jupyter Notebook and JupyterLab interfaces alongside an EMR console. An EMR notebook is a "serverless" notebook where the commands are executed using a kernel on the EMR cluster. The contents of the notebooks are also saved to Amazon S3 separately from cluster data for durability and flexible re-use. One can start a cluster, attach an EMR notebook for analysis, and then terminate the cluster. One can also close a notebook attached to one running cluster and switch to another. Multiple users can attach notebooks


```
print('Test set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
    test_loss, correct, len(data_loader.dataset),
    100. * correct / len(data_loader.dataset)))

In [43]:
if args.cuda:
    model = model.cuda()
    optimizer = optim.Adam(model.parameters(), lr=args.lr)
    for epoch in range(1, args.epochs + 1):

        train_epoch(epoch, args, model, train_loader, optimizer)
        test_epoch(model, test_loader)

Train Epoch: 1 [0/3339 (0%)]    Loss: 0.684990
Train Epoch: 1 [1280/3339 (38%)]    Loss: 0.463046
Train Epoch: 1 [2560/3339 (75%)]    Loss: 0.432949

Test set: Average loss: 4.6717, Accuracy: 465/590 (79%)


In [ ]:
```

B. AWS EMR through SSH in Name Node

[illegible]

VIII. CONCLUSION & FUTURE WORK

With this approach we were successful to accurately classify whether an MRI Scan has a tumor in it. We could take this further and use a segmentation model such as a ResUNet to determine and localize the tumor in the MRI Scan. We could also classify the type of tumor that exists in the MRI Scan to take it one level forward. When we don't have much data to train our model, we can use data augmentation to add to data



to get higher accuracy. We could also use data augmentation by making some changes in the existing dataset like flipping, rotating, scaling to achieve better results.

- [1] Dipanjan Sarkar. "A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning"
- [2] S. J. Pan and Q. Yang, "A Survey on Transfer Learning," in IEEE Transactions on Knowledge and Data Engineering
- [3] Ying Tai, Jian Yang, Xiaoming Liu; Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017
- [4] <https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-managed-notebooks.html>
- [5] Simard, Patrice Y., David Steinkraus, and John C. Platt. "Best practices for convolutional neural networks applied to visual document analysis." Icdar. Vol 3
- [6] Alom, Md. Zahangir Taha, Tarek Yakopcic, Christopher Westberg, Stefan Hasan, Mahmudul Esesn, Brian Awwal, Abdul Asari, Vijayan. (2018). The History Began from AlexNet: A Comprehensive Survey on Deep Learning Approaches.

- [1] Dipanjan Sarkar. "A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning"
- [2] S. J. Pan and Q. Yang, "A Survey on Transfer Learning," in IEEE Transactions on Knowledge and Data Engineering
- [3] Ying Tai, Jian Yang, Xiaoming Liu; Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017
- [4] <https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-managed-notebooks.html>
- [5] Simard, Patrice Y., David Steinkraus, and John C. Platt. "Best practices for convolutional neural networks applied to visual document analysis." Icdar. Vol 3
- [6] Alom, Md. Zahangir Taha, Tarek Yakopcic, Christopher Westberg, Stefan Hasan, Mahmudul Esesn, Brian Awwal, Abdul Asari, Vijayan. (2018). The History Began from AlexNet: A Comprehensive Survey on Deep Learning Approaches.